

System Programming

2nd Laboratory (8th and 11rd March 2017)

C programming revision (debug and valgrind)

In this laboratories all programs should be compiled with the options **-g -O0** (these options insert debug information into the programs and disable any optimization). Although some of the errors are evident and trivial do not correct the programs before following all the steps in the exercises.

Attaching processes to debugger

I

The program **infinite-loop.c** does not terminate. In order to understand where the infinite loop is, it is necessary to use the debugger.

I.a) The easiest way is to start the program inside the debugger:

- **ddd infinite-loop**
- press the button **run**
- after some time press the button **interrup**
- in the lower window observe the location
- in the lower window issue the command **where**

At this point it is possible to observe the point where the program has stopped.

If necessary issues the command **up** (several times) in the lower window.

When the debugger is in the upper frame (0x00000000004005af in main () at infinite-loop.c:10) it is possible to continue the program step by step:

- issue the command **display I**
- press **step** several times
- observe the value of **i**

I.b) if the program is running in the shell it is necessary to attach the debugger to the running process:

- start the program **inifnite-loop** in a terminal
- start a new terminal
- in the new terminal
 - issue the command **ps -a | grep inifilite-loop**
 - take note of the process ID (first number on the output of the command)
 - for example **658**
 - issue the command **ddd inifinite-loop**
 - in the lower window of **ddd**
 - issue the command **attach 658**

From this moment it is possible to use the debugger as if the application was started inside the debugger: (**where**, **up**, **print**, **display**, ...)

I.c) Correct the program.

DDD manual:

<https://www.gnu.org/software/ddd/manual/>

Core dump II

Sometimes the program crashes due to invalid pointers. When this happens the program stops running, a message is printed in the terminal and if configured a core file is generated:

- compile the program `char-conv.c`
- execute it in the terminal
- write a word and press enter

II.a) If the program is executed inside the debugger (**ddd**) from the beginning the program is interrupted and it is possible to observe the incorrect values:

- in the terminal run **ddd char-conv**
- inside **ddd**
 - run the program (press the button **run**)
 - write a word in the lower window
 - wait for the program to crash (a red arrow will show the wrong line)
 - issue the command **where**
 - issue the command **print v1**
 - issue the command **print v2**

The address of `v2` is invalid!

Exit **ddd**.

II.b) In order to do a postmortem evaluation of the program ran in the terminal it is necessary for the operating system to generate a core dump file. These commands work on ubuntu, may not work on other versions of linux :(

In the terminal:

- issue the command **ulimit -c unlimited**
- execute the application (**char-conv**)
- write a word
- the program crashes and a **core** file is generated
- issue the command **ls**
- execute the debugger issuing the command **ddd char-conv core**
- the debugger presents the location of the crash

Correct the error.

Overruns a leaks

III

The directory **III** contains a possible correction to the previous exercise (**char-conv-prob.c**). Although working correctly this program has two programming errors.

In order to find these issues **valgrind** can be used:

- compile the program
- run the program inside valgrid
- in the terminal: issue the command: **valgrind --leak-check=full -v ./char-conv-prob**
- observe the output

The last message of valgrind states that the program has 2 errors:

ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)

Memory leak

The first error described is a memory leak:

==2985== HEAP SUMMARY:

==2985== in use at exit: 7 bytes in 1 blocks

==2985== total heap usage: 3 allocs, 2 frees, 2,055 bytes allocated

this message informs that a malloc was made but no free was performed before the end of the program.

Correct this error.

Buffer overruns

Valgrind also identifies a **Invalid read of size 1**

This error means that during the program execution a read operations tried to access a memory outside a valid array:

- Address 0x51db8c7 is 0 bytes after a block of size 7 alloc'd

Valgrind also indicates that the memory was allocated in **main (char-conv-correct.c:15)** and the access was performed in **main (char-conv-correct.c:20)**

Correct this error.

Valgrind memcheck manual:

<http://valgrind.org/docs/manual/mc-manual.html>

Uninitialized memory

V

The directory **IV** contains a possible correction to the previous exercise (**char-conv-uninit.c**). Although working correctly this program still has one programming error.

In order to find this issues **valgrind** can be used:

- compile the program
- run the program inside valgrid
- in the terminal: issue the command: **valgrind --leak-check=full -v ./ char-conv-prob**
- observe the output

Valgrind states that a **Conditional jump or move depends on uninitialised value(s)** on line **main (char-conv-uninit.c:20)**. This happens because the conversion loop did not copy the '\0'

Correct the error.

Apps vs servers V

In the previous exercises (III and IV) we identified three different memory allocation related errors.

Describe why the programs managed to run correctly even in the presence of these errors:

--

For each of the previous errors describe how their presence in a server running continuously can affect the availability of the server:

Memory leak	
Buffer overrun	
Uninitialized memory	

Caches V

Compiles the programs inside directory **V**.

Run each of the programs issuing the following commands:

- `time ./cache-example1`
- `time ./cache-example2`

Why does **cache-example2** take more time to execute than **cache-example1**?

Valgrind also allows the evaluation of the cache usage of a program. Execute the following command in two different terminals:

- `valgrind --tool=cachegrind --cachegrind-out-file=cache-example1.out ./cache-example1`
- `valgrind --tool=cachegrind --cachegrind-out-file=cache-example2.out ./cache-example2`

Compare the two outputs and conclude why there is difference in execution time.

Valgrind also allows the presentation of the code lines that are less efficient (more cache misses). In two different terminals issue the following commands:

- `cg_annotate --auto=yes cache-example2.out`
- `cg_annotate --auto=yes cache-example1.out`

What is the line responsible for the performance difference?

Why is there this difference?

Valgrind Cachegrind manual:

<http://valgrind.org/docs/manual/cg-manual.html>

Overall performance VI

Valgrind also allows to evaluate the overall performance of the application. Allowing the comparison of the instructions/memory access/cache misses and predicted execution time of each function.

An extra program (**kcachegrind**) allows an easy study of the collected data.

To obtain the overall data, run:

- compile the program **complex.c**
- issue the command:
 - **valgrind --tool=callgrind --cache-sim=yes --callgrind-out-file=complex.out ./complex**
- wait for the simulation to end
- issue the command:
 - **kcachegrind complex.out**

Explore the application:

- Change from **Instruction fetch** to another counter (for example **Cycle Estimation**)
- Compare the column **Incl.** vs **Self**
- Select the **Call Graph**
- Compare the **extract_column** and **extract_row** functions
- Select **Callee Map**

What is the function that takes longer to execute?

Valgrind Callgrind manual:

<http://valgrind.org/docs/manual/ci-manual.html>